

# MACHINE LEARNING LAB-7:PCA&LDA

## Submitted by:

Name: **Stebin George**

Register Number: **21122061**

Class: **2MSCDS**

Time Taken: **5 hours**

## Lab Overview

### Objectives

**TO get to know more about the PCA and LDA two major dimensionality reduction techniques in ML**

### LIBRARIES:

In [33]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn import datasets
from sklearn.decomposition import PCA
```

### Questions:

- Part A. Perform PCA and LDA on Breast Cancer Dataset, write down your obsevations. While loading, use the toy dataset available in SKLearn (load\_breast\_cancer)
- Part B. Illustrate the effect of changing various method parameters of PCA and LDA. Compare the accuracies, and provide visualizations and interpretations for the evaluation metrics.
- Part C. "PCA could be used in applications such as Image Processing, to reduce the complexity of data and improve performance or to compress images". Justify this statement with your own findings.

### Problem Definition:

**The problem tries to give a introducrion about Pca and Lda in dimensionality reduction.**

### Approach

- Importing the necessary libraries
- analysing the data and doing the basic operations

- doing the EDA and pre-processing steps
  - Loading the PCA & LDA function
  - Plotting the accuracy
  - checking for the accuracy
  - usage of images to show compression using PCA
- 

## Sections

1. Libraries
  2. introduction
  3. EDA
  4. PCA&LDA
  5. Q&A
  6. Accuracy Table
  7. Conclusion
- 

## Loading Dataset:

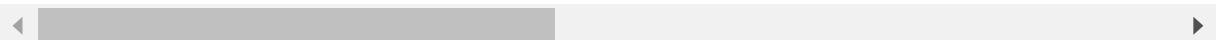
```
In [265...]: df=datasets.load_breast_cancer(as_frame=True)
df1=df.data
df1['diagnosis']=df.target
```

```
In [3]: df1.head()
```

Out[3]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	di
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	

5 rows × 31 columns



```
In [4]: #checking the shape
df1.shape
```

Out[4]: (569, 31)

**Datset has 569 rows and 33 columns**

```
In [5]: df1.describe()
```

Out[5]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension	mean radius error	mean texture error	mean perimeter error	mean area error	mean smoothness error	mean compactness error	mean concavity error	mean concave points error	mean symmetry error	mean fractal dimension error	mean worst radius	mean worst texture	mean worst perimeter	mean worst area	mean worst smoothness	mean worst compactness	mean worst concavity	mean worst concave points	mean worst symmetry	mean worst fractal dimension	mean diagnosis			
count	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000				
mean	14.127292	19.289649	91.969033	654.889104	0.096360	0.104341	0.088799	0.104341	0.096360	0.092630	0.092630	0.064920	0.064920	0.052813	0.052813	0.079720	0.079720	0.049560	0.049560	0.061540	0.061540	0.0345400	0.0345400	0.426800	0.426800	0.029560	0.029560	0.019380	0.019380	0.000000	0.000000	0.088799	0.088799	0.010000
std	3.524049	4.301036	24.298981	351.914129	0.014064	0.052630	0.052630	0.052630	0.014064	0.095870	0.095870	0.105300	0.105300	0.064920	0.064920	0.029560	0.029560	0.019380	0.019380	0.061540	0.061540	0.345400	0.345400	0.426800	0.426800	0.029560	0.029560	0.019380	0.019380	0.000000	0.000000	0.088799	0.088799	0.010000
min	6.981000	9.710000	43.790000	143.500000	0.052630	0.019380	0.019380	0.019380	0.052630	0.095870	0.095870	0.130400	0.130400	0.052813	0.052813	0.079720	0.079720	0.029560	0.029560	0.061540	0.061540	0.345400	0.345400	0.426800	0.426800	0.029560	0.029560	0.019380	0.019380	0.000000	0.000000	0.088799	0.088799	0.010000
25%	11.700000	16.170000	75.170000	420.300000	0.086370	0.064920	0.064920	0.064920	0.086370	0.105300	0.105300	0.130400	0.130400	0.052813	0.052813	0.079720	0.079720	0.029560	0.029560	0.061540	0.061540	0.345400	0.345400	0.426800	0.426800	0.029560	0.029560	0.019380	0.019380	0.000000	0.000000	0.088799	0.088799	0.010000
50%	13.370000	18.840000	86.240000	551.100000	0.095870	0.092630	0.092630	0.092630	0.095870	0.105300	0.105300	0.130400	0.130400	0.052813	0.052813	0.079720	0.079720	0.029560	0.029560	0.061540	0.061540	0.345400	0.345400	0.426800	0.426800	0.029560	0.029560	0.019380	0.019380	0.000000	0.000000	0.088799	0.088799	0.010000
75%	15.780000	21.800000	104.100000	782.700000	0.105300	0.130400	0.130400	0.130400	0.105300	0.163400	0.163400	0.345400	0.345400	0.052813	0.052813	0.079720	0.079720	0.029560	0.029560	0.061540	0.061540	0.426800	0.426800	0.029560	0.029560	0.019380	0.019380	0.000000	0.000000	0.088799	0.088799	0.010000		
max	28.110000	39.280000	188.500000	2501.000000	0.163400	0.345400	0.345400	0.345400	0.163400	0.426800	0.426800	0.029560	0.029560	0.019380	0.019380	0.000000	0.000000	0.088799	0.088799	0.010000	0.010000	0.029560	0.029560	0.019380	0.019380	0.000000	0.000000	0.088799	0.088799	0.010000				

8 rows × 31 columns

GIVES THE SUMMARY OF THE DATASET

## CHECKING NULL VALUES:

In [6]:

df1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   mean radius      569 non-null   float64
 1   mean texture     569 non-null   float64
 2   mean perimeter   569 non-null   float64
 3   mean area        569 non-null   float64
 4   mean smoothness  569 non-null   float64
 5   mean compactness 569 non-null   float64
 6   mean concavity   569 non-null   float64
 7   mean concave points 569 non-null   float64
 8   mean symmetry    569 non-null   float64
 9   mean fractal dimension 569 non-null   float64
 10  radius error     569 non-null   float64
 11  texture error    569 non-null   float64
 12  perimeter error  569 non-null   float64
 13  area error       569 non-null   float64
 14  smoothness error 569 non-null   float64
 15  compactness error 569 non-null   float64
 16  concavity error  569 non-null   float64
 17  concave points error 569 non-null   float64
 18  symmetry error   569 non-null   float64
 19  fractal dimension error 569 non-null   float64
 20  worst radius     569 non-null   float64
 21  worst texture    569 non-null   float64
 22  worst perimeter   569 non-null   float64
 23  worst area        569 non-null   float64
 24  worst smoothness 569 non-null   float64
 25  worst compactness 569 non-null   float64
 26  worst concavity   569 non-null   float64
 27  worst concave points 569 non-null   float64
 28  worst symmetry    569 non-null   float64
 29  worst fractal dimension 569 non-null   float64
 30  diagnosis        569 non-null   int32
```

```
dtypes: float64(30), int32(1)
memory usage: 135.7 KB
```

In [7]: `df1.isna().sum()`

```
Out[7]: mean radius          0
         mean texture         0
         mean perimeter        0
         mean area             0
         mean smoothness       0
         mean compactness      0
         mean concavity        0
         mean concave points   0
         mean symmetry          0
         mean fractal dimension 0
         radius error          0
         texture error          0
         perimeter error        0
         area error             0
         smoothness error       0
         compactness error      0
         concavity error        0
         concave points error   0
         symmetry error          0
         fractal dimension error 0
         worst radius           0
         worst texture          0
         worst perimeter         0
         worst area              0
         worst smoothness        0
         worst compactness       0
         worst concavity         0
         worst concave points   0
         worst symmetry          0
         worst fractal dimension 0
         diagnosis              0
         dtype: int64
```

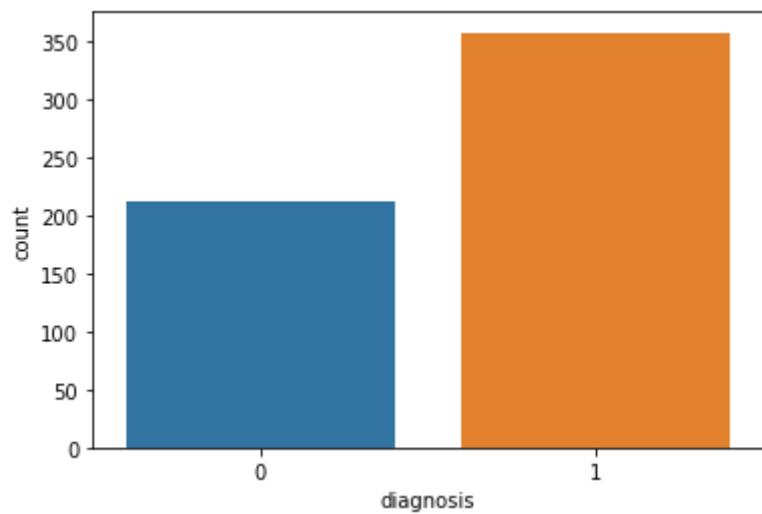
We dont have any null values here

## EDA:

In [8]: `import warnings  
warnings.filterwarnings('ignore')`

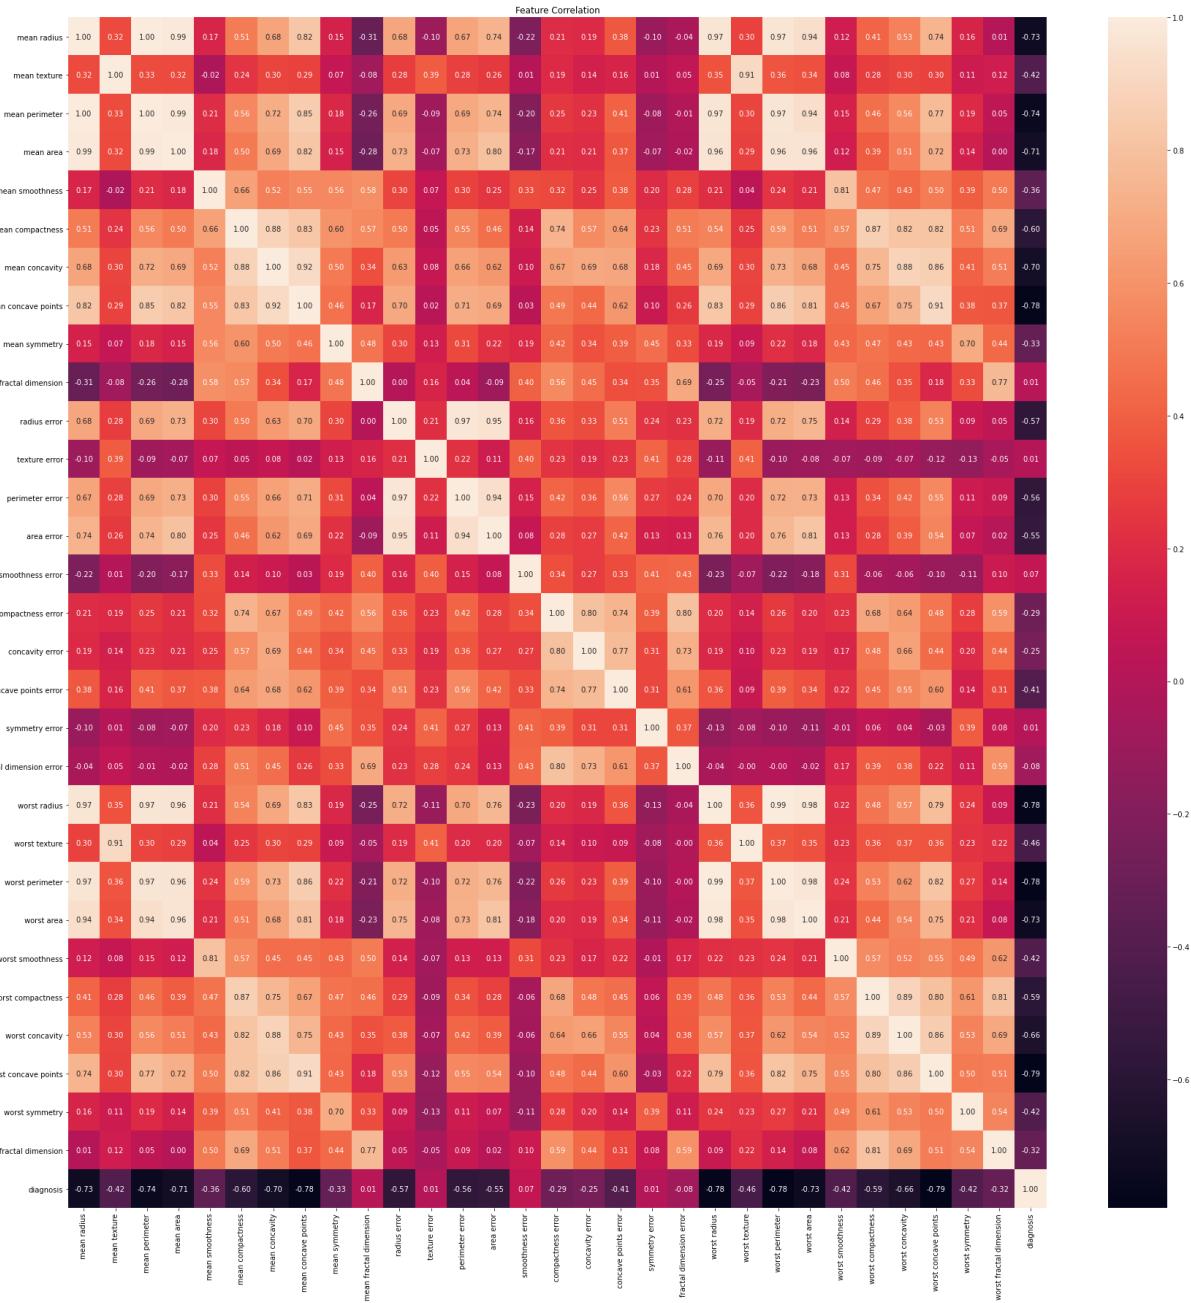
In [9]: `ax = sns.countplot(df1.diagnosis,label="Count")  
B, M =df1.diagnosis.value_counts()  
print('Number of Benign: ',B)  
print('Number of Malignant : ',M)`

```
Number of Benign:  357
Number of Malignant :  212
```



In [10]:

```
# Correlation Matrix
plt.figure(figsize=(30,30))
corr_matrix = df1.corr()
sns.heatmap(corr_matrix, annot = True, fmt = '.2f',)
plt.title("Feature Correlation")
plt.show()
```



The above graph shows the correlation plot of the different variables in the dataset.

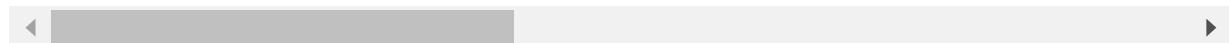
In [11]: `df1.corr()`

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	r con p
<b>mean radius</b>	1.000000	0.323782	0.997855	0.987357	0.170581	0.506124	0.676764	0.82
<b>mean texture</b>	0.323782	1.000000	0.329533	0.321086	-0.023389	0.236702	0.302418	0.29
<b>mean perimeter</b>	0.997855	0.329533	1.000000	0.986507	0.207278	0.556936	0.716136	0.85
<b>mean area</b>	0.987357	0.321086	0.986507	1.000000	0.177028	0.498502	0.685983	0.82
<b>mean smoothness</b>	0.170581	-0.023389	0.207278	0.177028	1.000000	0.659123	0.521984	0.55
<b>mean compactness</b>	0.506124	0.236702	0.556936	0.498502	0.659123	1.000000	0.883121	0.83

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	r con p
<b>mean concavity</b>	0.676764	0.302418	0.716136	0.685983	0.521984	0.883121	1.000000	0.92
<b>mean concave points</b>	0.822529	0.293464	0.850977	0.823269	0.553695	0.831135	0.921391	1.00
<b>mean symmetry</b>	0.147741	0.071401	0.183027	0.151293	0.557775	0.602641	0.500667	0.46
<b>mean fractal dimension</b>	-0.311631	-0.076437	-0.261477	-0.283110	0.584792	0.565369	0.336783	0.16
<b>radius error</b>	0.679090	0.275869	0.691765	0.732562	0.301467	0.497473	0.631925	0.69
<b>texture error</b>	-0.097317	0.386358	-0.086761	-0.066280	0.068406	0.046205	0.076218	0.02
<b>perimeter error</b>	0.674172	0.281673	0.693135	0.726628	0.296092	0.548905	0.660391	0.71
<b>area error</b>	0.735864	0.259845	0.744983	0.800086	0.246552	0.455653	0.617427	0.69
<b>smoothness error</b>	-0.222600	0.006614	-0.202694	-0.166777	0.332375	0.135299	0.098564	0.02
<b>compactness error</b>	0.206000	0.191975	0.250744	0.212583	0.318943	0.738722	0.670279	0.49
<b>concavity error</b>	0.194204	0.143293	0.228082	0.207660	0.248396	0.570517	0.691270	0.43
<b>concave points error</b>	0.376169	0.163851	0.407217	0.372320	0.380676	0.642262	0.683260	0.61
<b>symmetry error</b>	-0.104321	0.009127	-0.081629	-0.072497	0.200774	0.229977	0.178009	0.09
<b>fractal dimension error</b>	-0.042641	0.054458	-0.005523	-0.019887	0.283607	0.507318	0.449301	0.25
<b>worst radius</b>	0.969539	0.352573	0.969476	0.962746	0.213120	0.535315	0.688236	0.83
<b>worst texture</b>	0.297008	0.912045	0.303038	0.287489	0.036072	0.248133	0.299879	0.29
<b>worst perimeter</b>	0.965137	0.358040	0.970387	0.959120	0.238853	0.590210	0.729565	0.85
<b>worst area</b>	0.941082	0.343546	0.941550	0.959213	0.206718	0.509604	0.675987	0.80
<b>worst smoothness</b>	0.119616	0.077503	0.150549	0.123523	0.805324	0.565541	0.448822	0.45
<b>worst compactness</b>	0.413463	0.277830	0.455774	0.390410	0.472468	0.865809	0.754968	0.66
<b>worst concavity</b>	0.526911	0.301025	0.563879	0.512606	0.434926	0.816275	0.884103	0.75
<b>worst concave points</b>	0.744214	0.295316	0.771241	0.722017	0.503053	0.815573	0.861323	0.91
<b>worst symmetry</b>	0.163953	0.105008	0.189115	0.143570	0.394309	0.510223	0.409464	0.37

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	r con p
worst fractal dimension	0.007066	0.119205	0.051019	0.003738	0.499316	0.687382	0.514930	0.36
diagnosis	-0.730029	-0.415185	-0.742636	-0.708984	-0.358560	-0.596534	-0.696360	-0.77

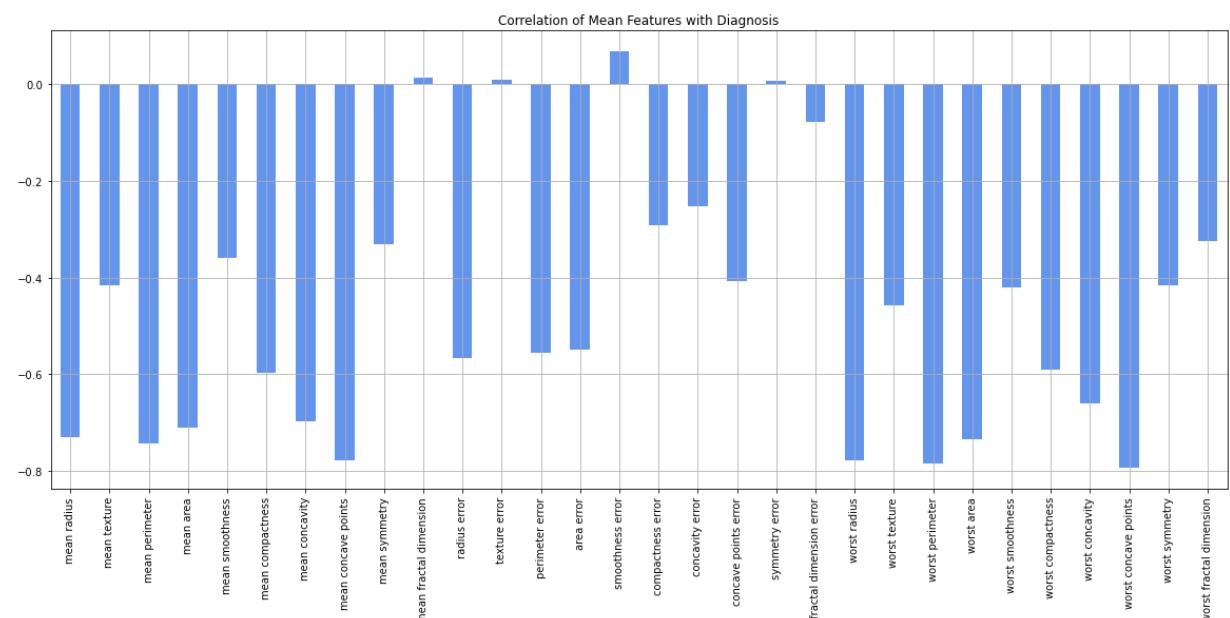
31 rows × 31 columns



## Feature contributions to the target variable:

In [19]:

```
df_mean = df1[df1.columns[:]]
plt.figure(figsize=(20, 8))
df_mean.drop('diagnosis', axis=1).corrwith(df_mean.diagnosis).plot(kind='bar', grid=
```



In [28]:

```
X = df1.iloc[:, 1:-1]
y = df1.iloc[:, -1]
```

## Standardizing the dataset:

In [110...]

```
sc = StandardScaler()
X_stand = sc.fit_transform(X)
```

## Question1

Part A. Perform PCA and LDA on Breast Cancer Dataset, write down your obsevations. While loading, use the toy dataset available in SKLearn (load\_breast\_cancer)

## NORMAL DATA:

In [97]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.8, random_state
```

```
In [98]: classifier = LogisticRegression()
          classifier.fit(X_train, y_train)
          predictions = classifier.predict(X_test)
          predictions
```

```
In [99]: print(accuracy_score(y_test,predictions))
```

0.9407894736842105

## PCA:

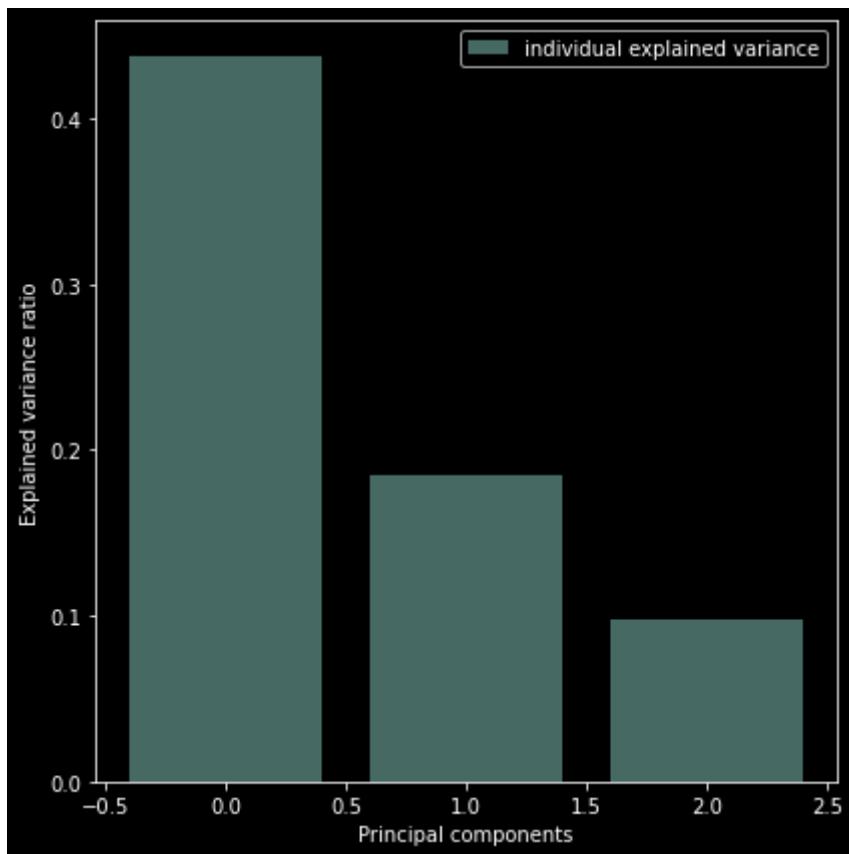
```
In [275...]:  
pca = PCA(n_components=3)  
pca.fit(X_stand)  
PCA(copy=True, iterated_power='auto', n_components=3, random_state=None,  
    svd_solver='auto', tol=0.0, whiten=False)  
X_pca = pca.transform(X_stand)  
X_pca.shape
```

Out[275... (569, 3)

```
In [276]: explained_variance=pca.explained_variance_ratio_
          explained variance
```

```
Out[276]: array([0.43706363, 0.18472237, 0.09716239])
```

```
In [277...]:  
with plt.style.context('dark_background'):  
    plt.figure(figsize=(6,6))  
  
    plt.bar(range(3), explained_variance, alpha=0.5, align='center',  
            label='individual explained variance')  
    plt.ylabel('Explained variance ratio')  
    plt.xlabel('Principal components')  
    plt.legend(loc='best')  
    plt.tight_layout()
```



```
In [114...]: X_train, X_test, y_train, y_test = train_test_split(X_pca, y, test_size=0.8, random_s
```

```
In [115...]: classifier = LogisticRegression()
          classifier.fit(X_train, y_train)
```

```
Out[115... LogisticRegression()
```

```
In [116]: predictions = classifier.predict(X_test)  
predictions
```

```
In [117...]: print(accuracy_score(y_test, predictions))
```

0.956140350877193

## OBSERVATION:

With PCA the logistics regression was giving better accuracy PCA has no concern with the class labels. It summarizes the feature set without considering the output. PCA tries to find the directions of the maximum variance in the dataset. In a high cardinality feature set, there are possibilities of duplicate features which would add redundancy to the dataset, increase the computation cost and add unnecessary model complexity. The role of PCA is to find such highly correlated or duplicate features and to come up with a new feature set where there is minimum correlation between the features or in other words feature set with maximum variance between the features.

## LDA:

```
In [118...]: X_train, X_test, y_train, y_test = train_test_split(X_stand, y, test_size=0.8, random
```

```
In [119]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA  
  
lda = LDA(n_components = 1)  
X_train = lda.fit_transform(X_train, y_train)  
X_test = lda.transform(X_test)
```

```
In [120...]: classifier = LogisticRegression()
          classifier.fit(X_train, y_train)
```

```
Out[120]: LogisticRegression()
```

```
In [121...]: predictions = classifier.predict(X_test)  
predictions
```

```
In [122]: print(accuracy_score(y_test, predictions))
```

0.9320175438596491

## OBSERVATIONS:

LDA tries to reduce the dimensionality by taking into consideration the information that discriminates the output classes. LDA tries to find the decision boundary around each cluster of class. It projects the data points to new dimension in a way that the clusters are as separate from each other as possible and individual elements within a class are as close to the centroid as possible. In other words, the inter-class separability is increased in LDA. Intra-class separability is reduced. The new dimensions are the linear discriminants of the feature set.

## QUESTION-2:

**Part B. Illustrate the effect of changing various method parameters of PCA and LDA. Compare the accuracies, and provide visualizations and interpretations for the evaluation metrics.**

In [258...]

```
def standard_scalar(data):
    sc = StandardScaler()
    X_stand = sc.fit_transform(data)
    return(X_stand)

def pca(X_stand,n_components):
    pca = PCA(n_components)
    pca.fit(X_stand)
    X_pca = pca.transform(X_stand)
    return (X_pca)

def logistic(X_train, X_test, y_train, y_test):
    classifier = LogisticRegression()
    classifier.fit(X_train, y_train)
    predictions = classifier.predict(X_test)
    acc=accuracy_score(y_test,predictions)
    return(acc)
```

In [261...]

```
method=[ 'NORMAL' , 'PCA' , 'LDA' ]
randomstate=[5,10,15,20,25]
n_components=[1,3,5,7,9]
methods=[]
accu=[]
random=[]
n_component=[]
for i in method:
    if i =='NORMAL':
        for j in randomstate:
            methods.append(i)
            random.append(j)
            n_component.append('NA')
            X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.8,
            acc=logistic(X_train, X_test, y_train, y_test)
            accu.append(acc)
    elif i=='PCA':
        X_stand=standard_scalar(X)
        for j in randomstate:
            for k in n_components:
                methods.append(i)
                random.append(j)
                n_component.append(k)
                X_pca=pca(X_stand,k)
```

```

        X_train, X_test, y_train, y_test = train_test_split(X_pca, y, test_size=1)
        acc1=logistic(X_train, X_test, y_train, y_test)
        accu.append(acc1)

    elif i=='LDA':
        for j in randomstate:
            methods.append(i)
            random.append(j)
            n_component.append(1)
            X_train, X_test, y_train, y_test = train_test_split(X_stand, y, test_size=1)
            lda = LDA(n_components = 1)
            X_train = lda.fit_transform(X_train, y_train)
            X_test = lda.transform(X_test)
            acc1=logistic(X_train, X_test, y_train, y_test)
            accu.append(acc1)

```

## ACCURACY DATAFRAME:

In [262]...

```

score=pd.DataFrame()
score['method']=methods
score['randomstate']=random
score['no of components']=n_component
score['accuracy']=accu
score

```

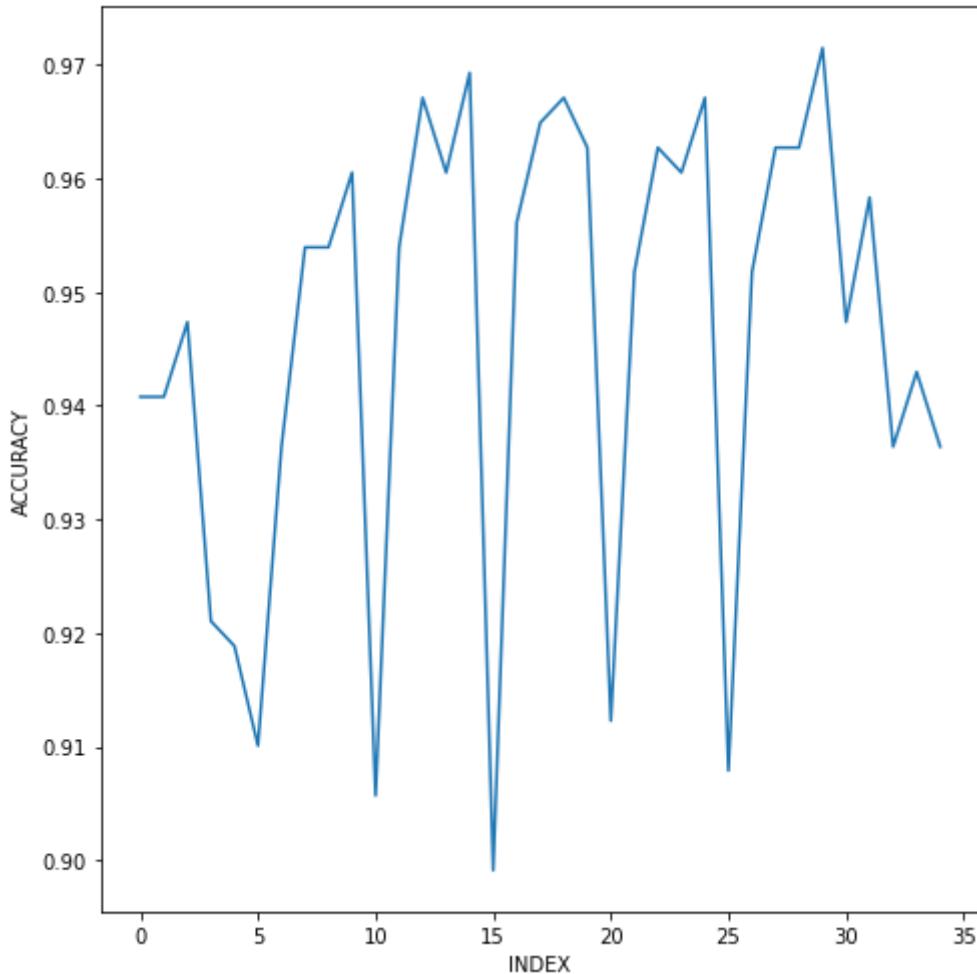
Out[262]...

	method	randomstate	no of components	accuracy
0	NORMAL	5	NA	0.940789
1	NORMAL	10	NA	0.940789
2	NORMAL	15	NA	0.947368
3	NORMAL	20	NA	0.921053
4	NORMAL	25	NA	0.918860
5	PCA	5	1	0.910088
6	PCA	5	3	0.936404
7	PCA	5	5	0.953947
8	PCA	5	7	0.953947
9	PCA	5	9	0.960526
10	PCA	10	1	0.905702
11	PCA	10	3	0.953947
12	PCA	10	5	0.967105
13	PCA	10	7	0.960526
14	PCA	10	9	0.969298
15	PCA	15	1	0.899123
16	PCA	15	3	0.956140
17	PCA	15	5	0.964912
18	PCA	15	7	0.967105
19	PCA	15	9	0.962719
20	PCA	20	1	0.912281

	method	randomstate	no of components	accuracy
<b>21</b>	PCA	20	3	0.951754
<b>22</b>	PCA	20	5	0.962719
<b>23</b>	PCA	20	7	0.960526
<b>24</b>	PCA	20	9	0.967105
<b>25</b>	PCA	25	1	0.907895
<b>26</b>	PCA	25	3	0.951754
<b>27</b>	PCA	25	5	0.962719
<b>28</b>	PCA	25	7	0.962719
<b>29</b>	PCA	25	9	0.971491
<b>30</b>	LDA	5	1	0.947368
<b>31</b>	LDA	10	1	0.958333
<b>32</b>	LDA	15	1	0.936404
<b>33</b>	LDA	20	1	0.942982
<b>34</b>	LDA	25	1	0.936404

In [282...]

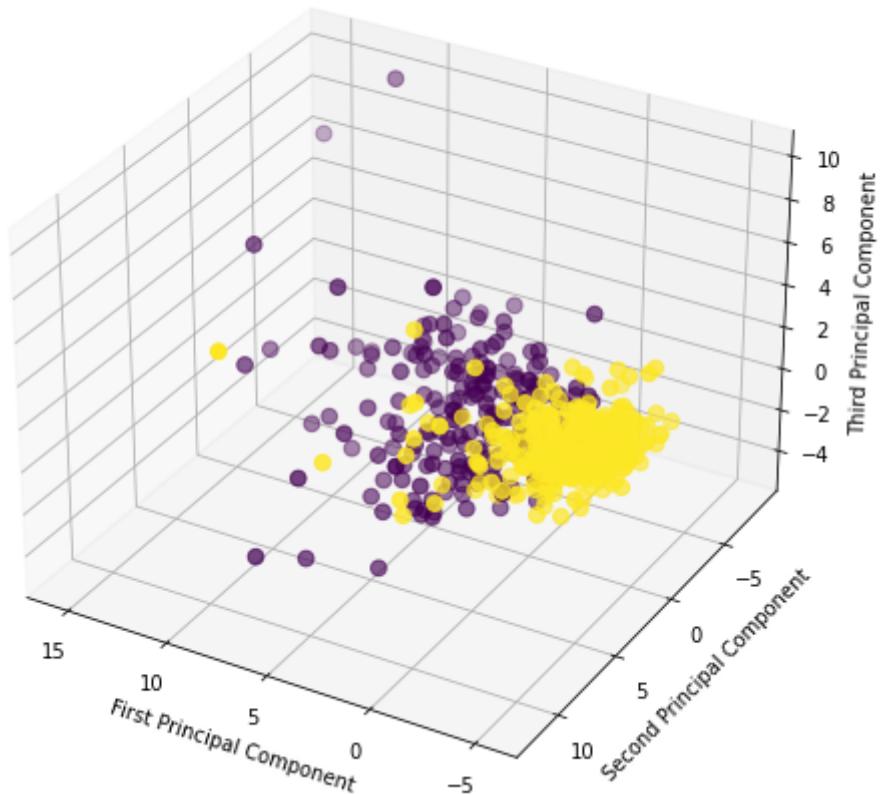
```
plt.figure(figsize=(7,7))
plt.plot(accu)
plt.ylabel('ACCURACY')
plt.xlabel('INDEX')
plt.tight_layout()
```



For the random state of 25 and a pca component of 9 we get the maximum accuracy of 97.14%

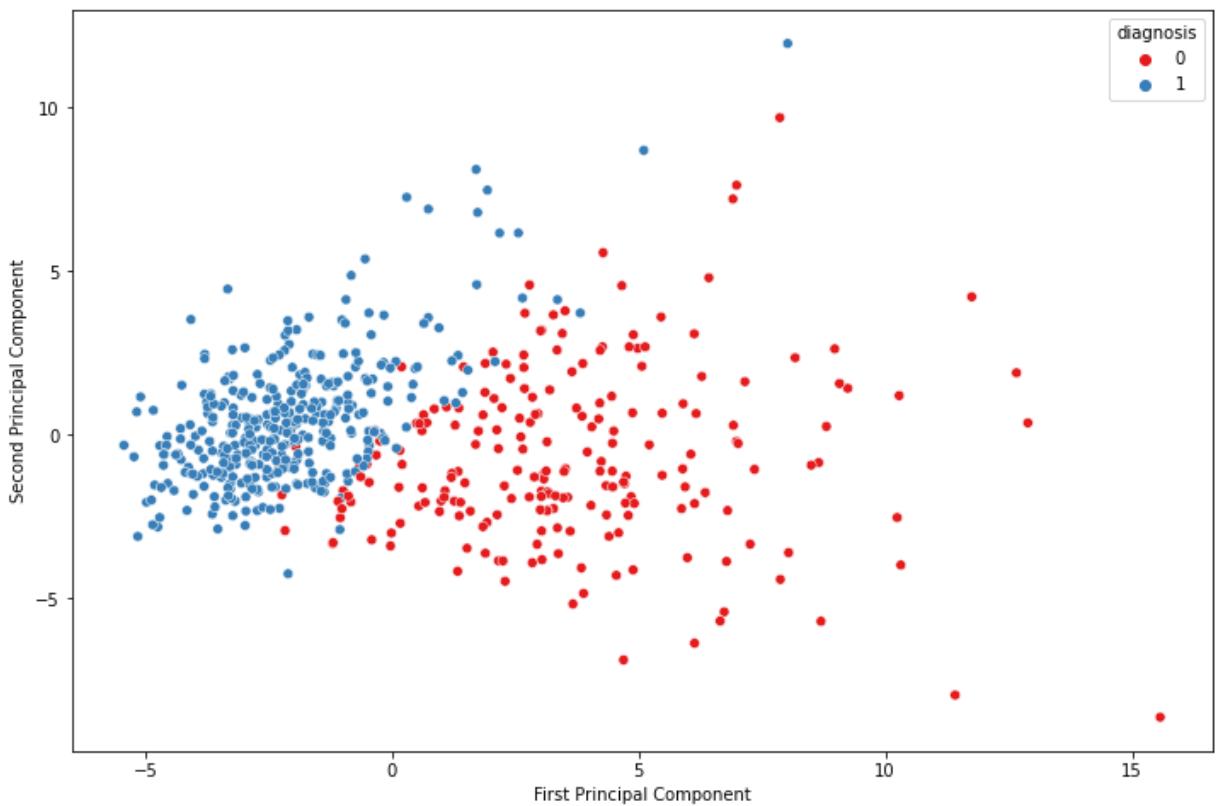
In [263]:

```
fig = plt.figure(figsize=(15, 8))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x_pca[:,0], x_pca[:,1], x_pca[:,2], c=df1['diagnosis'], s=60)
ax.set_xlabel('First Principal Component')
ax.set_ylabel('Second Principal Component')
ax.set_zlabel('Third Principal Component')
ax.view_init(30, 120)
```



```
In [264...]  
ax = plt.figure(figsize=(12,8))  
sns.scatterplot(x_pca[:,0], x_pca[:,1],hue=df1['diagnosis'], palette ='Set1' )  
plt.xlabel('First Principal Component')  
plt.ylabel('Second Principal Component')
```

Out[264... Text(0, 0.5, 'Second Principal Component')



## QUESTION 3:

"PCA could be used in applications such as Image Processing, to reduce the complexity of data and improve performance or to compress images". Justify this statement with your own findings.

The MNIST dataset contains the image data of handwritten digits.

In [239...]

```
mnist = pd.read_csv(r"C:\Users\stebi\OneDrive\Desktop\mnist.csv")
mnist.head()
```

Out[239...]

	label	1x1	1x2	1x3	1x4	1x5	1x6	1x7	1x8	1x9	...	28x19	28x20	28x21	28x22	28x23
0	5	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
2	4	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
3	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
4	9	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0

5 rows × 785 columns

The 'label' column contains the values of the digit (0–9). We do not need that column for our analysis because PCA is an unsupervised machine learning task that does not deal with labelled data. So, we can simply drop that column.

In [241...]

```
mnist.drop(columns='label', inplace=True)
mnist.head()
```

Out[241...]

	1x1	1x2	1x3	1x4	1x5	1x6	1x7	1x8	1x9	1x10	...	28x19	28x20	28x21	28x22	28x23
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0

5 rows × 784 columns

In [240...]

```
mnist.shape
```

Out[240...]

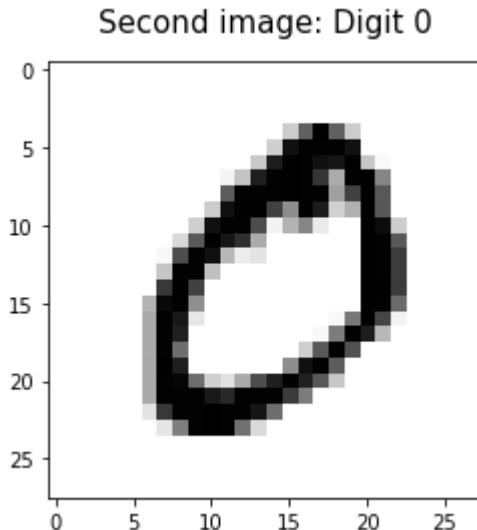
(60000, 785)

### Plotting the image:

In [246...]

```
second_image = mnist.iloc[1].values.reshape([28,28])
plt.imshow(second_image, cmap='gray_r')
```

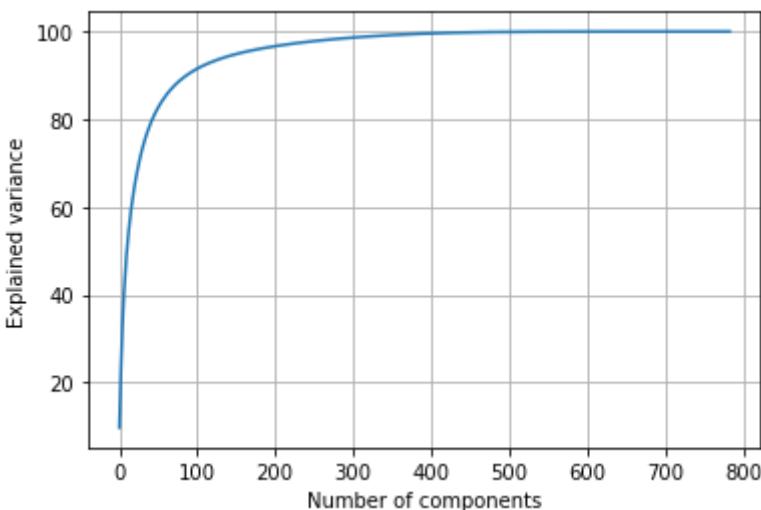
```
plt.title('Second image: Digit 0', fontsize=15, pad=15)
plt.savefig("Second image.png")
```



**Apply PCA with the original number of dimensions (i.e., 784) and create the scree plot to see how well PCA captures the variance of the data.**

```
In [255...]
pca_784 = PCA(n_components=784)
pca_784.fit(mnist)
plt.grid()
plt.plot(np.cumsum(pca_784.explained_variance_ratio_* 100))
plt.xlabel('Number of components')
plt.ylabel('Explained variance')
```

Out[255...]



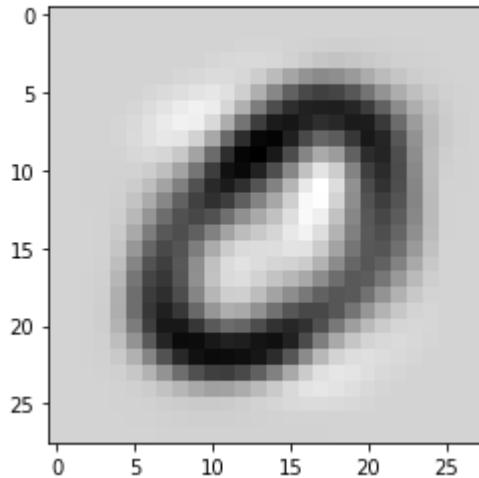
**Reducing the number of principal components to 10 and obtaining the output**

```
In [248...]
pca_10 = PCA(n_components=10)
mnist_pca_10_reduced = pca_10.fit_transform(mnist)
mnist_pca_10_recovered = pca_10.inverse_transform(mnist_pca_10_reduced)

image_pca_10 = mnist_pca_10_recovered[1,:].reshape([28,28])
plt.imshow(image_pca_10, cmap='gray_r')
plt.title('Compressed image with 10 components', fontsize=15, pad=15)
```

Out[248...]

### Compressed image with 10 components



**Increasing the number of components of pca to 184 shows much more clearer image than pca with 10 components.**

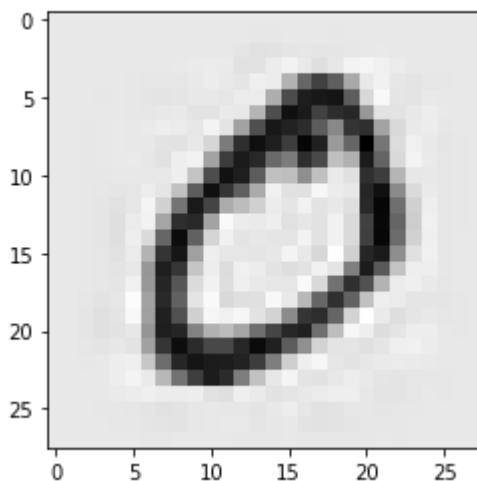
In [253...]

```
pca_184 = PCA(n_components=184)
mnist_pca_184_reduced = pca_184.fit_transform(mnist)
mnist_pca_184_recovered = pca_184.inverse_transform(mnist_pca_184_reduced)

image_pca_184 = mnist_pca_184_recovered[1,:].reshape([28,28])
plt.imshow(image_pca_184, cmap='gray_r')
plt.title('Compressed image with 184 components', fontsize=15, pad=15)
```

Out[253...]: Text(0.5, 1.0, 'Compressed image with 184 components')

### Compressed image with 184 components



In [254...]

```
np.cumsum(pca_184.explained_variance_ratio_ * 100)[-1]
```

Out[254...]: 96.11902137200546

**96.12% of the variance in the model is explained by the pca with 184 instead of 784. Thus we actually reduced the size of the image without actually losing the accuracy of the model.**

**OBSERVATION:**

The image at the left is the original image with 784 dimensions. The image at the right is the compressed image with 184 dimensions. After applying PCA on image data, the dimensionality has been reduced by 600 dimensions while keeping about 96% of the variability in the original image data! By comparing these two images, you can see that there is a slight image quality loss, but the content of the compressed image is still visible.

---

## References:

PCA

<https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html>  
<https://towardsdatascience.com/principal-component-analysis-for-breast-cancer-data-with-r-and-python-b312d28e911f> <https://www.kaggle.com/jahirmorenoa/pca-to-the-breast-cancer-data-set> [https://www.youtube.com/watch?v=e2sM7ccaA9c&ab\\_channel=DigitalSreeni](https://www.youtube.com/watch?v=e2sM7ccaA9c&ab_channel=DigitalSreeni)  
<https://www.datacamp.com/community/tutorials/principal-component-analysis-in-python>  
<https://towardsdatascience.com/dimensionality-reduction-of-a-color-photo-splitting-into-rgb-channels-using-pca-algorithm-in-python-ba01580a1118>  
<https://www.kaggle.com/mirzarahim/introduction-to-pca-image-compression-example>  
[https://github.com/gtraskas/breast\\_cancer\\_prediction/blob/master/breast\\_cancer.ipynb](https://github.com/gtraskas/breast_cancer_prediction/blob/master/breast_cancer.ipynb)

LDA

[http://scikit-learn.org/stable/modules/generated/sklearn.discriminant\\_analysis.LinearDiscriminantAnalysis.html](http://scikit-learn.org/stable/modules/generated/sklearn.discriminant_analysis.LinearDiscriminantAnalysis.html)  
<https://machinelearningmastery.com/linear-discriminant-analysis-with-python/>  
<https://towardsdatascience.com/linear-discriminant-analysis-in-python-76b8b17817c2>  
<https://www.mygreatlearning.com/blog/linear-discriminant-analysis-or-lda/>  
<https://www.geeksforgeeks.org/ml-linear-discriminant-analysis/>

In [ ]: